

Million_Song_Subset

February 9, 2025

1 Data Science in Music: A Machine Learning Recommendation Model

1.1 Overview

This notebook implements a **music recommendation system** using machine learning and data-driven techniques. The goal is to recommend **relevant songs** based on a user's music preferences, leveraging **data science** and **AI-powered algorithms**.

1.2 Key Features

- **Data Preprocessing:** Cleans and structures music-related data.
- **Recommendation Engine:** Uses similarity algorithms to suggest songs.
- **Visualization & Insights:** Analyzes patterns in the dataset.

1.3 Dataset

The dataset contains various song attributes such as **artist names**, **genres**, **track features**, and **popularity metrics**.

1.4 Approach

1. Data Cleaning & Preprocessing
2. Feature Engineering & Similarity Calculation
3. Building the Recommendation Model
4. Exploratory Data Analysis (EDA)

Let's dive into the implementation!

```
[2]: # Import necessary libraries
import os
import h5py
import pandas as pd
import numpy as np
from glob import glob
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
```

```
import seaborn as sns
from sklearn.decomposition import PCA
import boto3
```

```
[13]: # Path to one of your HDF5 files (modify this path based on your dataset
      ↪structure)
file_path = "A:\Data science\MillionSongSubset\A\A\A\TRAAAW128F429D538.h5"

def explore_hdf5_file(file_path):
    """Load and explore an HDF5 file from the Million Song Dataset"""
    with h5py.File(file_path, "r") as h5_file:
        print(" Keys in this file:", list(h5_file.keys())) # Check available
        ↪groups

        # Explore the metadata group
        if "metadata" in h5_file:
            metadata = h5_file["metadata"]
            print("\n Artist Name:", metadata["songs"]["artist_name"][0].
            ↪decode("utf-8"))
            print(" Artist ID:", metadata["songs"]["artist_id"][0].
            ↪decode("utf-8"))
            print(" Title:", metadata["songs"]["title"][0].decode("utf-8"))
            #print(" Release Year:", metadata["songs"]["year"][0])

        # Explore the analysis group
        if "analysis" in h5_file:
            analysis = h5_file["analysis"]
            print("\n Tempo:", analysis["songs"]["tempo"][0])
            print(" Key:", analysis["songs"]["key"][0])
            print(" Duration:", analysis["songs"]["duration"][0])

        # Explore the musicbrainz group
        if "musicbrainz" in h5_file:
            musicbrainz = h5_file["musicbrainz"]
            #print("\n Album Name:", musicbrainz["songs"]["release"][0].
            ↪decode("utf-8"))

# Run exploration function
explore_hdf5_file(file_path)
```

Keys in this file: ['analysis', 'metadata', 'musicbrainz']

Artist Name: Casual

Artist ID: ARD7TVE1187B99BFB1

Title: I Didn't Mean To

Tempo: 92.198

Key: 1

Duration: 218.93179

<>:2: SyntaxWarning: invalid escape sequence '\D'

<>:2: SyntaxWarning: invalid escape sequence '\D'

C:\Users\ASUS\AppData\Local\Temp\ipykernel_30484\618296351.py:2: SyntaxWarning: invalid escape sequence '\D'

file_path = "A:\Data science\MillionSongSubset\A\A\A\TRAAAW128F429D538.h5"

```
[ ]: # Define the columns we want (optional)
selected_features = [
    "artist_name", "artist_hotttnesss", "title", "tempo", "key", "duration",
    "loudness"
]

def extract_features(file_path):
    """Extract selected features from a single HDF5 file."""
    with h5py.File(file_path, "r") as f:
        metadata = f["metadata"]["songs"]
        analysis = f["analysis"]["songs"]

        data = {}
        for col in selected_features:
            if col in metadata.dtype.names:
                value = metadata[col][0]
            elif col in analysis.dtype.names:
                value = analysis[col][0]
            else:
                value = None # If the column does not exist

            # Handle different data types
            if isinstance(value, np.bytes_): # If byte string, decode
                data[col] = value.decode("utf-8") if value else "Unknown"
            elif isinstance(value, (int, float, np.integer, np.floating)): # Numeric
                data[col] = float(value)
            else:
                data[col] = "Unknown" # Default for missing values

            # Handle artist terms separately (since it's a list)
            if "artist_terms" in metadata.dtype.names:
                artist_terms = metadata["artist_terms"][0]
                data["artist_terms"] = [x.decode('utf-8') if isinstance(x, bytes)
                else x for x in artist_terms]
            else:
                data["artist_terms"] = []

        return data
```

```

# Example: Load a single file
file_path = "A:/Data science/MillionSongSubset/A/A/A/TRAAAW128F429D538.h5"
df_song = extract_features(file_path)

# Convert to DataFrame
df_sample = pd.DataFrame([df_song])
df_sample.head()

```

```

[ ]:  artist_name  artist_hotttnesss      title  tempo  key  duration \
0      Casual          0.401998  I Didn't Mean To  92.198  1.0  218.93179

      loudness artist_terms
0    -11.197          []

```

```

[ ]: # Base directory where all HDF5 files are stored
base_dir = "A:/Data science/MillionSongSubset"

# Find all HDF5 files recursively
all_files = glob(os.path.join(base_dir, "**/*.h5"), recursive=True)

# Process all files and store them in a DataFrame
data_list = []
for file in all_files:
    try:
        song_data = extract_features(file)
        data_list.append(song_data)
    except Exception as e:
        print(f"Error processing {file}: {e}")

# Convert collected data into a single DataFrame
df_msd = pd.DataFrame(data_list)

# Save to CSV for future use
df_msd.to_csv("million_song_dataset.csv", index=False)

```

```

[ ]: # Apply StandardScaler
test_df = df_msd.copy()
scaler = StandardScaler()
test_df[features] = scaler.fit_transform(test_df[features])

```

```

[ ]: # Normalize numerical Data
# Select only numerical features for clustering
features = ["tempo", "key", "duration", "loudness", "artist_hotttnesss"]

# Apply StandardScaler
test_df = df_msd.copy()

```

```
scaler = StandardScaler()
test_df[features] = scaler.fit_transform(test_df[features])
```

```
[19]: # Train the Recommendation Model
# Train the K-Means Model
kmeans = KMeans(n_clusters=5, random_state=42)
test_df["cluster"] = kmeans.fit_predict(test_df[features])

test_df[["title", "artist_name", "cluster"]]
```

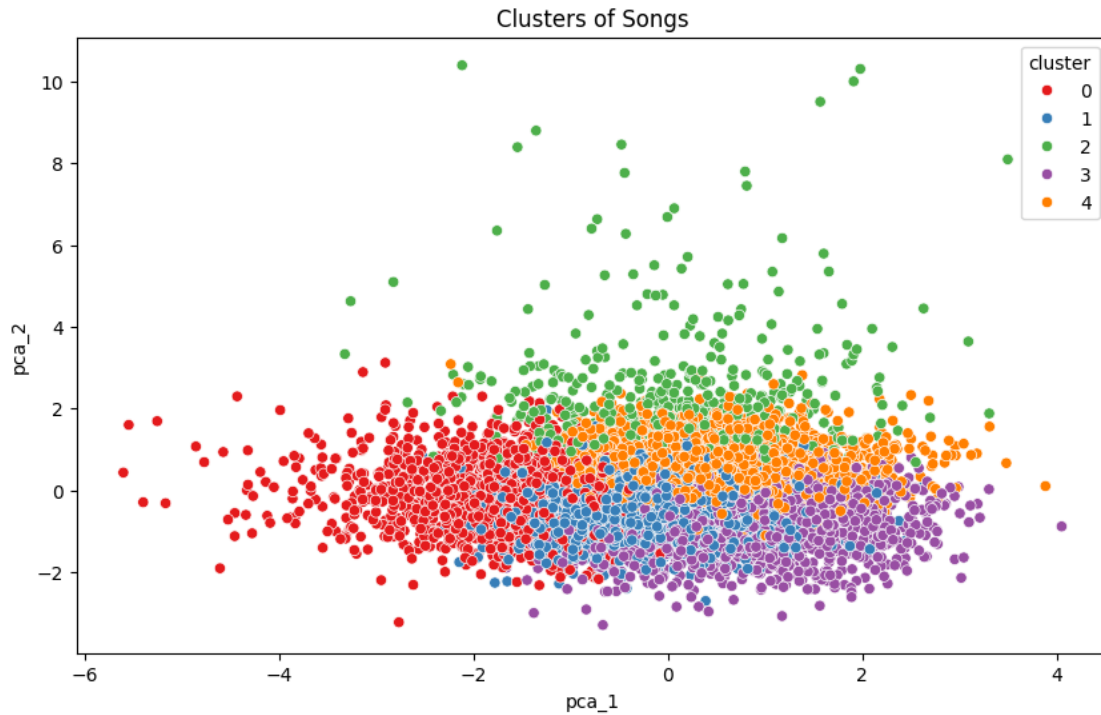
```
[19]:
```

	title	artist_name	cluster
0	I Didn't Mean To	Casual	1
1	Soul Deep	The Box Tops	4
2	Amor De Cabaret	Sonora Santanera	4
3	Something Girls	Adam Ant	1
4	Face the Ashes	Gob	1
...
9995	The Hanged Man	Moonspell	4
9996	The Wonderful World Of The Young	Danny Williams	0
9997	Sentimental Man	Winston Reedy	1
9998	Zydeco In D-Minor	Myrick "Freeze" Guillory	1
9999	Shattered Life	Seventh Day Slumber	1

[10000 rows x 3 columns]

```
[20]: # Reduce data to 2D using PCA
pca = PCA(n_components=2)
test_df["pca_1"], test_df["pca_2"] = pca.fit_transform(test_df[features])[:, 0], \
    pca.fit_transform(test_df[features])[:, 1]

# Plot clusters
plt.figure(figsize=(10, 6))
sns.scatterplot(x=test_df["pca_1"], y=test_df["pca_2"], hue=test_df["cluster"], \
    palette="Set1")
plt.title("Clusters of Songs")
plt.show()
```



```
[21]: # Load the dataset
# Ensuring the DF is ready to save as csv for uploading in supabase
test_df.insert(0, "ID", range(0, len(test_df)))

DB = pd.DataFrame()

DB["track_id"] = test_df["ID"].astype(str) # Primary Key
DB["artist_name"] = test_df["artist_name"].astype(str)
DB["track_name"] = test_df["title"].astype(str)
DB["cluster_id"] = test_df["cluster"].astype(str)

# Save proper dataset for Supabase upload
DB.to_csv("MSD_clustered.csv", index=False)
```

```
[ ]: # Generate Music Recommendations
# Test recommendation

# Function to find recommended artists
def recommend_artists(artist_name):
    # Check if artist exists
    if artist_name not in DB["artist_name"].values:
        return f" Artist '{artist_name}' not found in the dataset."

    # Get the cluster of the given artist
```

```

artist_cluster = DB[DB["artist_name"] == artist_name]["cluster_id"].iloc[0]

# Find other artists in the same cluster
recommended_artists = DB[DB["cluster_id"] == artist_cluster]["artist_name"].
↪unique()

# Remove the input artist from recommendations
recommended_artists = [artist for artist in recommended_artists if artist !
↪= artist_name]

return {
    "input_artist": artist_name,
    "cluster_id": int(artist_cluster),
    "recommended_artists": recommended_artists[:5] # Show max 10
↪recommendations
}

# Test the function
artist_input = input("Enter an artist name: ").strip()
result = recommend_artists(artist_input)
print(result)

```

```

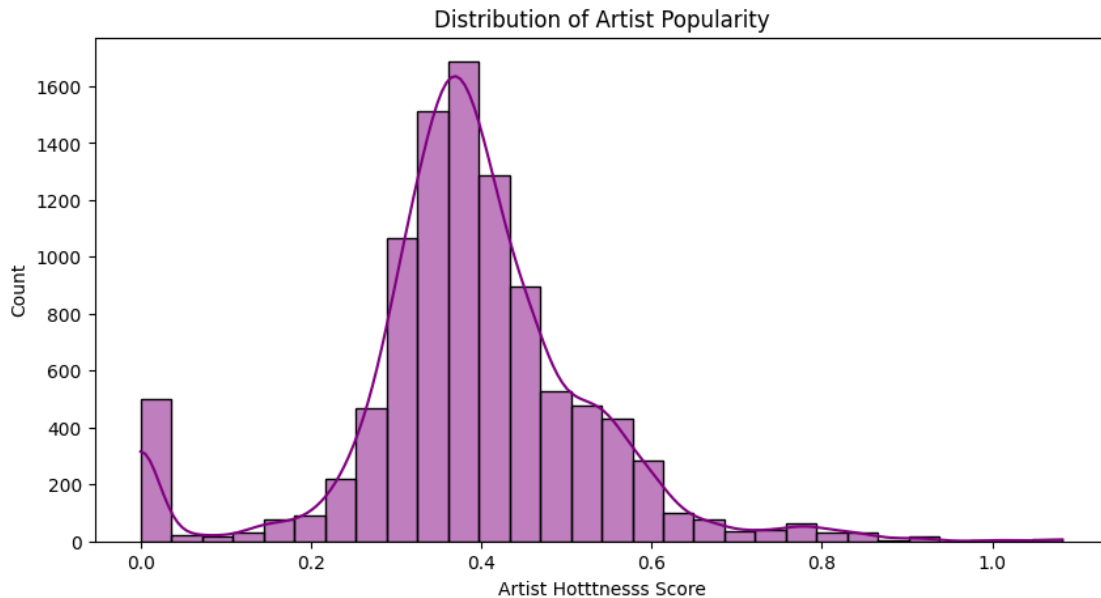
{'input_artist': 'Eminem', 'cluster_id': 1, 'recommended_artists': ['Casual',
'Adam Ant', 'Gob', 'Tweeterfriendly Music', 'Lionel Richie']}

```

```

[28]: # Histogram of artist hotness
plt.figure(figsize=(10, 5))
sns.histplot(df_msd['artist_hotttnesss'], bins=30, kde=True, color="purple")
plt.title("Distribution of Artist Popularity")
plt.xlabel("Artist Hotttnesss Score")
plt.ylabel("Count")
plt.show()

```



```
[70]: # Top 10 artists by popularity
top_artists = df_msd.groupby("artist_name")["artist_hotttnesss"].mean().
    ↪sort_values(ascending=False).head(10)

plt.figure(figsize=(12, 6))

sns.barplot(x=top_artists.index, y=top_artists.values, palette="magma")

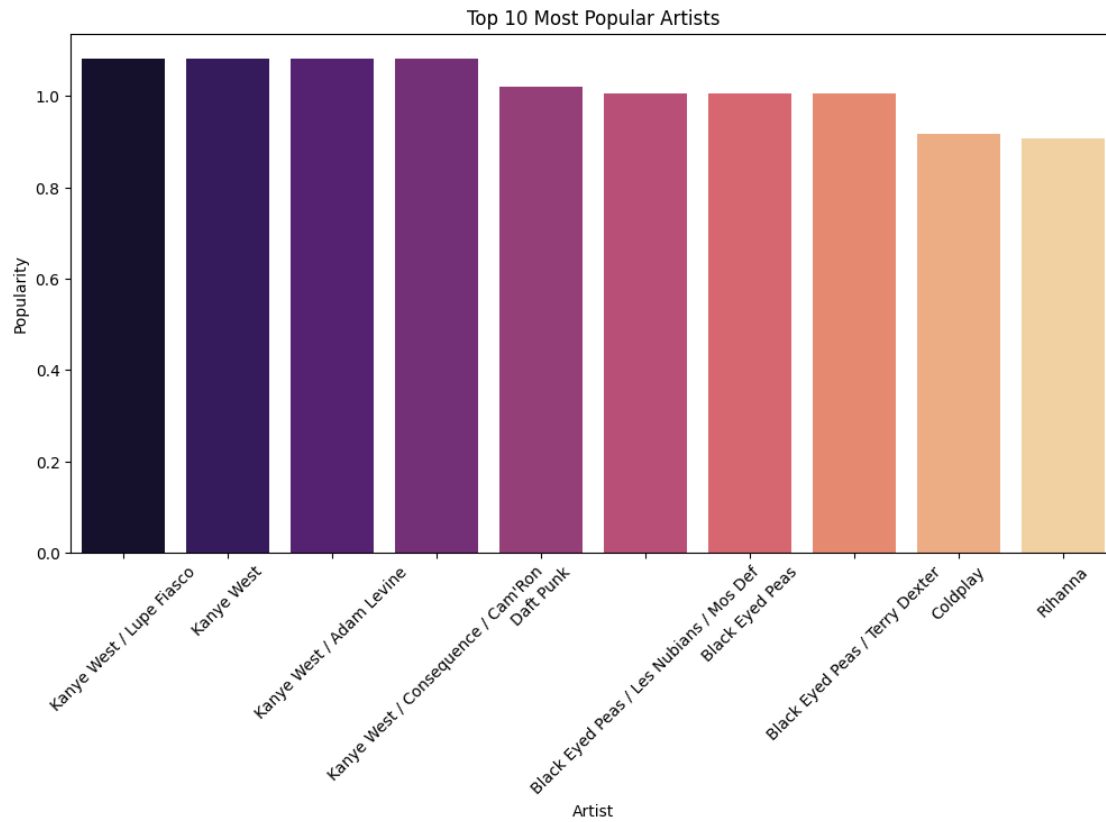
plt.title("Top 10 Most Popular Artists")
plt.xticks(rotation=45)
plt.xlabel("Artist")
plt.ylabel("Popularity")

plt.show()
```

C:\Users\ASUS\AppData\Local\Temp\ipykernel_30484\3347214049.py:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

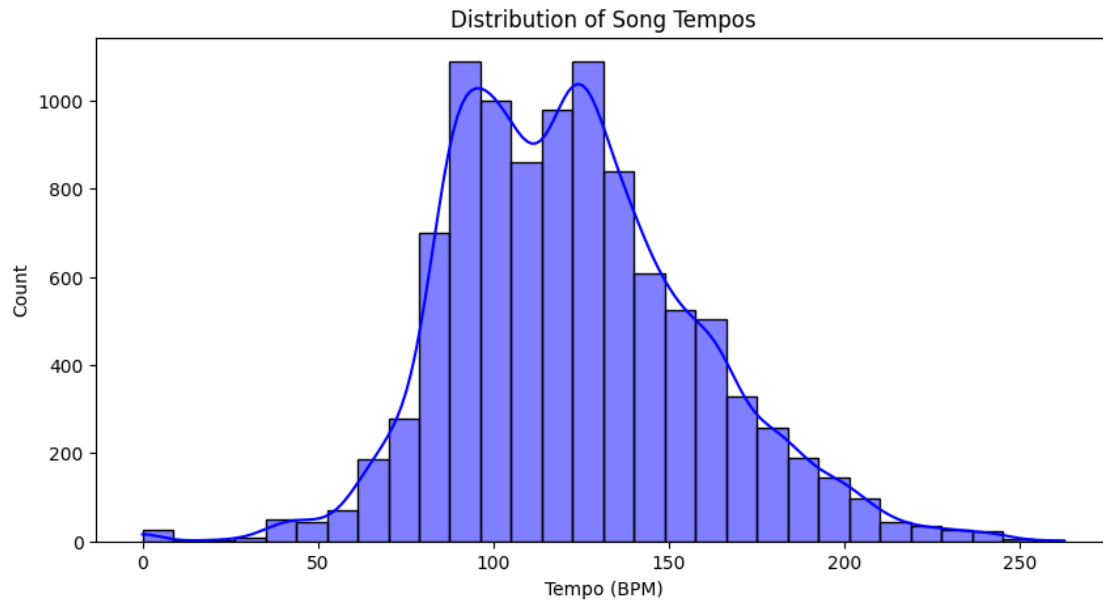
```
sns.barplot(x=top_artists.index, y=top_artists.values, palette="magma")
```

```
[69]: # Histogram of song tempos
plt.figure(figsize=(10, 5))

sns.histplot(df_msd['tempo'], bins=30, kde=True, color="blue")

plt.title("Distribution of Song Tempos")
plt.xlabel("Tempo (BPM)")
plt.ylabel("Count")
plt.show()
```



```
[68]: plt.figure(figsize=(12, 6))

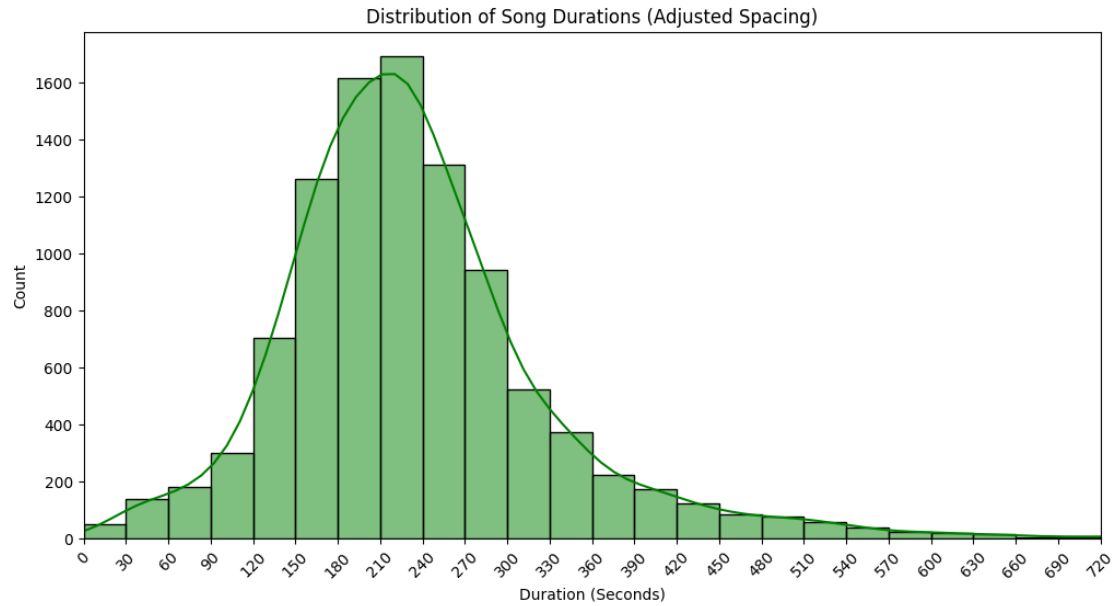
# Define bins for better spacing (every 30 seconds up to the max duration)
bins = np.arange(0, df_msd["duration"].max() + 30, 30)
sns.histplot(df_msd["duration"], bins=bins, kde=True, color="green")

# Increase spacing on x-axis by setting custom limits
plt.xlim(0, 330)

# Set major xticks at 30-second intervals
plt.xticks(np.arange(0, 750, 30), rotation=45)

plt.title("Distribution of Song Durations (Adjusted Spacing)")
plt.xlabel("Duration (Seconds)")
plt.ylabel("Count")

plt.show()
```



```
[76]: # Observe the correlation btw. tempo & loudness

import seaborn as sns
plt.figure(figsize=(10, 6))
sns.scatterplot(x=df_msd['tempo'], y=df_msd['loudness'], alpha=0.6, color="red")
plt.title('Tempo vs. Loudness')
plt.xlabel('Tempo')
plt.ylabel('Loudness')
plt.show()
```

